

7. Combinational Logic

Armed with the abstract model of combinational devices outlined in Chapter 5 and the concrete implementation technology for simple gates of Chapter 6, we turn our attention to techniques for constructing combinational circuits that perform arbitrarily complex useful functions. To this end, we use the constructive property of combinational devices outlined in Section 5.3.2: an acyclic circuit whose components are combinational devices is itself a combinational device. That is, the behavior of such a circuit can be described by the combination of a functional and timing specification derived from the circuit topology and the component specifications.

Functional Specifications


7.1. Functional Specifications

There are a variety of notations we can use to specify the functional specification of a combinational device such as the 3-input, single-output module shown to the right. We might simply use a sentence or algorithmic description of each output such as that shown in the box. A more systematic specification is the truth tables that have been used in prior examples; an equivalent truth table specification for this example is shown to the left. A concise alternative to the truth table is a Boolean expression:

$$Y = \bar{C}\bar{B}A + \bar{C}BA + C\bar{B}\bar{A} + CBA$$

represents the same 3-input function as the truth table and informal algorithmic specification. In each case, the combinational device abstraction requires that the functional description specify a logical (0 or 1) value for each combination of input values.

Boolean Expressions



7.2. Boolean Expressions

Boolean expressions are a ubiquitous notation in the engineering of digital systems. They constitute a simple algebra for functions defined on the domain of truth values (0 and 1), and define three basic operations:

- Multiplication (logical AND): $A \cdot B$ or simply AB represents the logical conjunction of the values A and B : the result is 1 only if A and B are both 1.
- Addition (logical OR): $A + B$ represents the logical disjunction of the values A and B ; the result is 1 if $A = 1$ or $B = 1$ or both.
- Negation (logical NOT or inversion): \bar{A} represents the logical inverse of A : $\bar{A} = 1$ if $A = 0$, $\bar{A} = 0$ if $A = 1$.

The use of familiar algebraic notation -- multiplication and addition -- for logical AND and OR operations takes a bit of time for the uninitiated to assimilate. The following list summarizes some properties of Boolean expressions that are worth internalizing:

Boolean Identities

OR rules:	$a + 1 = 1$	$a + 0 = a$
	$a + a = a$	
AND rules:	$a \cdot 1 = a$	$a \cdot 0 = 0$
	$a \cdot a = a$	
Commutativity:	$a + b = b + a$	$a \cdot b = b \cdot a$
Associativity:	$(a + b) + c = a + (b + c)$	
	$(a \cdot b) \cdot c = a \cdot (b \cdot c)$	
Distributivity:	$a \cdot (b + c) = a \cdot b + a \cdot c$	
	$a + b \cdot c = (a + b) \cdot (a + c)$	

Complements:	$\overline{a + a} = 1$	$\overline{a \cdot a} = 0$
	$\overline{\overline{a}} = a$	
Absorption:	$a + a \cdot b = a$	$a + \overline{a} \cdot b = a + b$
	$\overline{a \cdot (a + b)} = a$	$\overline{a \cdot (a + b)} = ab$
Reduction:	$\overline{a \cdot b + a \cdot b} = b$	$\overline{(a + b) \cdot (a + b)} = b$
DeMorgan's Law:	$\overline{\overline{a + b}} = a \cdot b$	$\overline{\overline{a \cdot b}} = a + b$

As the identities in the above table shows, there are many equivalent ways of representing a given logic function as a Boolean expression, and a variety of techniques for simplifying a Boolean expression or converting it to a form that is convenient for one's purposes.

While there are many Boolean expressions that represent a given Boolean function, there is essentially only one truth table for each function: modulo permutation of inputs, a truth table can be viewed as a canonical representation. One can verify the equivalence of two Boolean expressions with k variables by building truth tables that enumerate their values for each of the 2^k combinations of variable values, and checking that the tables are identical

Truth Table Arithmetic

7.3. Truth Table Arithmetic

The truth table for a k -input, single-output logic function has 2^k rows, listing the single-bit output for each of the 2^k potential combinations of input values. Each such k -input truth table has the identical form: input values are listed, conventionally in a systematic order corresponding to binary counting order; truth tables for different k -input functions thus differ only in the 2^k single-bit values in their output column.

There is a one-to-one correspondence between these truth tables and the set of k -input logic functions. As there are 2^k output bits that distinguish a particular k -input function, and each output bit may be 0 or 1, it follows that the total number of possible k -input functions is just 2^{2^k} -- precisely the number of distinct k -input truth tables.

There are, for example, a total of $2^{2^2} = 16$ 2-input functions, whose truth tables are summarized below:

Note that the full set of 2-input functions includes degenerate cases of functions whose output is independent of one or both input values. It is worth noting that, of these sixteen 2-input

functions, only six can be implemented directly as single CMOS gates: $NOR(A, B) = \overline{A + B}$ and $NAND(A, B) = \overline{A \cdot B}$, as well as the degnerate cases A , B , 0, and 1. Other 2-input functions must be synthesized using multiple CMOS gates in a combinational circuit.

Gates and Gate Symbols

7.4. Gates and Gate Symbols

Circuit symbols for commonly-used 2-input combinational logic components are shown to the right; some of these (NAND and NOR) we've already seen in connection with their direct implementation as CMOS gates. The others require multiple CMOS gates to implement. The 2-input AND of A and B can be easily implmented as a NAND gate followed by an inverter, since $A \cdot B = \overline{\overline{A \cdot B}} = \overline{NAND(A, B)}$; similarly, OR can be implemented as a CMOS

NOR gate followed by an inverter. The more complex $XOR(A,B) = A \cdot B + A \cdot \overline{B}$ requires several CMOS gates to implement, as does its inverse $XNOR(A,B)$.

Note the consistent use of small circles -- "inversion bubbles" -- to indicate inversion of logic levels.

Universal Gates

7.5. Universal Gates

Although there are 2^k k -input logic functions we may wish to implement as combinational devices, the ability to synthesize new devices as acyclic circuits implies that we can begin with a small repertoire of basic gates and design circuits using them to build arbitrary functionality. Indeed, the observation that Boolean expressions are built using only three operations (AND, OR, and NOT) implies that a starting toolkit containing only AND, OR, and inverter modules is sufficient. Moreover, since we can synthesize k -input AND and OR devices from circuits of 2-input gates, it follows that an unlimited supply of 2-input AND and OR devices, along with inverters, will allow us to synthesize any logic functions we might wish. We term such a set universal, in that it contains all we need to implement arbitrary logic functions.

Happily, the 2-input NAND and NOR operations that are easily implemented as single CMOS gates are each, in isolation, universal.

Given the universality of the *AND, OR, NOT* set, We can demonstrate the universality of *NAND* or *NOR* by showing how to implement each of *AND, OR, NOT* using only NAND or NOR gates, as shown to the right. From a logical standpoint, this implies that arbitrarily complex combinational systems can be built from a parts inventory containing only 2-input NAND or NOR gates. Sum-of-Products Synthesis

NAND, NOR Universality

7.6. Sum-of-Products Synthesis

It is straightforward to convert a truth table to an equivalent Boolean expression in sum-of-products form. Such an expression takes the form of the Boolean sum (OR) of several terms, each of which is the product of some set of (possibly inverted) input variables. Each such term t_i corresponds to a line of the truth table for which the output column contains 1 -- called an implicant of the function, since $t_i=1$ implies that the output is 1.

If we reconsider the prior 3-input example, we note that the four 1s in the output column imply that the output Y is 1 for each of the input combinations $ABC = 001, 011, 110, 111$. For each of these, we construct an implicant term that is 1 for precisely the specified combination of input values, by including each input variable in the term and inverting it (via the overbar notation) if its value is 0. The resulting sum-of-products expression is $\overline{C}\overline{B}A + \overline{C}BA + C\overline{B}\overline{A} + CBA$, an expression whose value is 1 for precisely those combinations of input values for which the truth table specifies $Y=1$.

SOP Implementation

Given that Boolean expressions are built from using AND, OR, and NOT (inversion) operations, it is tempting to synthesize circuits from Boolean expressions directly using AND gates, OR gates, and inverters, as shown to the left. Note that this approach generalizes to a systematic 3-tier structure for arbitrary k -input functions:

- a layer of inverters to generate the necessary inverted inputs;
- A set of N k -input AND-gates, where N is the number of implicants. Each AND gate computes the value of one term (implicant) of the sum-of-products expression; and
- An N -input OR gate to sum the terms.

Although the direct sum-of-products implementation seems a straightforward approach to implementation of arbitrary Boolean functions, it is rarely the fastest or most efficient (in

terms, say, of transistor count) of the possible CMOS implementation techniques. This stems primarily from two factors:

1. CMOS gates are naturally inverting: AND and OR gates use more transistors than NAND and NOR gates; optimized CMOS implementations exploit the simpler inverting logic for circuit simplicity, which is often at odds with the conceptual simplicity of Boolean expressions built from AND/OR/NOT operations.
2. Many Boolean functions have internal structure that can be exploited to reduce implementation complexity, even using direct AND/OR/NOT implementations. Such structure can often be uncovered by manipulation of a Boolean expression to yield a simpler, equivalent expression.

The number of transistors used in the above SOP implementation gives a rough proxy for its implementation cost. Our example uses four inverters (at two transistors each), four 3-input AND gates each requiring 8 transistors, and one 4-input OR requiring 10 transistors. The total transistor count for this circuit is thus 50; we will see shortly that the same function can be implemented far more cheaply.

Boolean Simplification

7.6.1. Boolean Simplification

Although a direct sum-of-products representation for a Boolean function is easy to derive systematically from a truth table, it rarely represents the function's simplest or most convenient form. We often can simplify the sum-of-product expressions by algebraic manipulation, using the identities of Section 7.2 and other tools.

Consider, for example, the unsimplified Boolean expression for the 3-input function of the previous sections:

$$Y = \bar{C}\bar{B}A + \bar{C}BA + C\bar{B}\bar{A} + CBA$$

Although this structure of the SOP implementation shown above corresponds directly to the structure of this circuit, we can find simpler expressions that suggest correspondingly simpler implementations of the same function by algebraic manipulation using the identities given above. A particularly useful identity is the reduction identity $a \cdot b + a \cdot \bar{b}$, which generalizes to:

$$aX\beta + a\bar{X}\beta = a\beta$$

where a and β are arbitrary (possibly empty) Boolean subexpressions. Thus we can simplify a Boolean expression by looking for pairs of terms that are identical except for some variable X , which is inverted in one term but not in the other, and by replacing the pair of terms by a single simpler term which omits the variable that omits X .

Applying this reduction step to the 3-input Boolean expression above, we notice that the terms $C\bar{B}\bar{A}$ and CBA match the pattern and can be replaced by the simpler term CB , yielding

$$Y = \bar{C}\bar{B}A + CB + \bar{C}BA$$

Noticing that $\bar{C}\bar{B}A$ and $\bar{C}BA$ can be reduced to $\bar{C}A$, we simply further to

$$Y = \bar{C}A + CB$$

which is a simpler sum-of-products Boolean equation that is equivalent to our original 4-term unreduced equation. In fact it has two implicants rather than four and each implicant is simpler, Simplified SOP circuit involving two variables rather than three. Each of the terms of this expression implies a 1 for two input combinations, allowing us to use fewer implicants in our simplified expression. In fact each term of the simplified expression is a prime implicant, meaning that no it is not subsumed by any other (yet simpler) implicant for the function.

In addition to suggesting simpler implementations, the simplified expression makes the function easier to understand: quick inspection reveals that $Y = A$ when $C = 0$, but $Y = B$ when $C = 1$. Although this characterization applies equally to the equivalent 4-term unsimplified expression, that relationship is much less obvious.

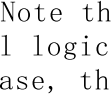
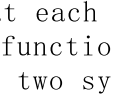
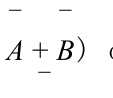

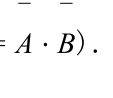

Based on this simplified Boolean expression, the direct sum-of-products implementation shown to the right is correspondingly less complex than that of our unsimplified 4-term SOP expression. This implementation uses a single inverter (2 transistors), two 2-input ANDs (6 transistors each) and a single 2-input OR (6 transistors) for a total transistor count of 20.

Repeated application of reduction steps together with other basic identities (such as reordering variables in a term to recognise reduction opportunities) is the approach taken by most software tools for Boolean simplification.

Pushing Bubbles

7.6.2. Pushing Bubbles

Another of the Boolean identities listed in Section 7.2 is often referred to as DeMorgan's Law, and constitutes a useful tool for transformation of Boolean expressions to more useful forms. One statement of Demorgan's law is $A + B = A \cdot B = NAND(A, B)$: that is, an OR with inverted inputs computes precisely the same value as an AND with an inverted output. Conversely, it follows that $NOR(A, B) = A + B = A \cdot B$.

Demorganized gate symbols	It is sometimes useful to reflect these equivalences in logic circuits by using "DeMorganized versions of circuit symbols, such as those shown to the left.	
		
	Note that each row of the table shows two different circuit symbols that perform the identical logic function -- indeed, they represent precisely the same circuit implementation. In each case, the two symbols reflect an application of DeMorgan's law:	
		
	inverting each input and output of an AND operation to convert it to an OR (since $A \cdot B = A + B$) or	
		
	inverting each input and output of an OR operation to convert it to an AND (since $A + B = A \cdot B$).	

By judicious use of these equivalences we often can replace conceptual AND and OR operations by the cheaper and faster NAND and NOR operations preferred in our CMOS technology. The use of alternative symbols (together with our understanding that pairs of inversion bubbles on the same wire cancel each other) can make such transformations easier to recognise in our circuit diagrams.

The diagram to the right shows a NAND/NAND implementation of the simplified sum-of-products expression that has been our running example. Relative to the AND/OR implementation, we have replace AND gates by NANDs (saving 2 transistors each) and used a Demorganized NAND gate instead of an OR to add the terms (saving another 2 transistors). The total transistor count of this implementation is 14. Of course, it computes the same value $Y = \bar{C}A + CB$ as our original 50-transistor circuit.

Circuit Timing

7.7. Circuit Timing

As noted in Section 6.4.4, we can bound the timing specifications t_{PD} and t_{CD} of an acyclic circuit of combinational components by considering cumulative delays along each input-to-output path through the circuit. To illustrate, consider our 3-input NAND-NAND circuit implemented using CMOS gates whose timing specifications are summarized to the left. We may bound the propagation delay by the longest cumulative t_{PD} along any input-output path; in this case, the constraining path involves an inverter and two NAND gates for a total propagation delay of 8ns. Similarly, we bound the contamination delay by the shortest cumulative t_{CD} along any input-output path, in this case involving 2 NAND gates for a total contamination delay of 2ns.		
Device	t_{PD}	t_{CD}
Inverter	2ns	1ns
NAND	3ns	1ns
Circuit	8ns	2ns

To explore the detailed timing of signals flowing through the circuit as the result of an input change, consider an experiment where the A and B inputs of our NAND-NAND Circuit Signals

circuit are held to 1 for an interval that includes a $1 \rightarrow 0$ transition on the C input. The diagram to the right includes new labels P , Q , and R attached to internal nodes of the circuit. The diagram to the left details the timing of each signal in the circuit as we maintain $A = B = 1$ and apply an instantaneous $1 \rightarrow 0$ transition on C at time $t = 0$. Note that, prior to the C transition, $A = B = C = 1$ and all signals in the circuit are at stable, valid values; the output $Y = 1$, as dictated by the functional specification of the circuit. The $1 \rightarrow 0$ transition on C at $t = 0$ stimulates the propagation of signal changes through the circuit, eventually affecting the output.

It is worth diving into the details of the timing of this signal flow to understand the guarantees provided by the (lenient) CMOS gates in our circuit:

1. The inverter output P remains 0 for 1ns, due to the $t_{CD} = 1ns$ inverter specification. It becomes a valid 1 at $t = 2$ -- 2ns following the input transition -- as guaranteed by the inverter's $t_{PD} = 2$ specification. Note that during the intervening 1ns interval, the signal on P is unspecified -- it may be a valid 0, a valid 1, or invalid. This interval of undetermined logic value is shown on the diagram by a crosshatched X pattern.
2. The output R of the lower NAND gate remains 0 for the 1ns NAND t_{CD} , and becomes a valid 1 after the 3ns NAND t_{PD} , with an intervening 2ns period of unspecified (possibly invalid) value.
3. The output Q of the upper NAND gate doesn't see an input change until $t = 1$, when its P input becomes invalid. Q remains a valid 1 for an additional 1ns contamination delay, after which it is contaminated. Q becomes a valid 0 at $t = 5$, the earliest time at which its inputs have been valid for the 3ns NAND propagation delay.
4. The output gate sees R go invalid at $t = 1$, but maintains a valid $Y = 1$ output for its 1ns contamination delay. Its R input becomes a valid 1 at $t = 3$, but its Q input remains invalid until becoming 0 at $t = 5$. Since a single 1 input is not sufficient to determine the output of even a lenient NAND gate, the output Y cannot be assumed to be valid until the 3ns propagation delay after both inputs become valid. Thus, $Y = 1$ at $t = 8$, following a 6ns interval of unspecified output value.

It is reassuring that, as the timing diagram shows, the output Y remains valid for 2ns following the input transition and becomes valid 8ns after that transition; this is consistent with the circuit's t_{CD} and t_{PD} bounds computed earlier.

It may be bothersome, however, that the output of our circuit is unspecified for an interval between the two periods of valid 1s. Given that $A = B = 1$, the transition on C does not affect the value of $Y = \bar{C}A + CB$, and we might ask that the output remain a valid 1 during this transition.

During the unspecified period, however, the output of our non-lenient device may become a valid 0, or even vacillate between 0 and 1 several times, before settling back to its final value of 1. Non-lenient behavior
Such anomalies are often referred to as glitches or hazards.

The fact that the use of lenient components does not guarantee lenience of acyclic circuits built from them is annoying, but usually not problematic. In fact this device is a valid combinational device (for which any input change contaminates the output value for a propagation delay) but not a lenient combinational device (whose output remains valid despite changes in irrelevant inputs).

For reasons explored in the next Chapter, reliable digital systems require a timing discipline that makes signal values significant only during prescribed intervals, allowing transient periods of invalidity during which they may make transitions. A simple such discipline is described in Section 8.3.3.

Lenient Circuits

7.7.1. Lenient Circuits

In certain important applications, however, lenience of a combinational component is critical. Suppose we had an application that requires a lenient device that computes $Y = \bar{C}A + CB$? We know that single CMOS gates are lenient, but a bit of analysis convinces us that this function cannot

be implemented as a single CMOS gate; hence we must resort to a circuit of CMOS components. Fortunately, it is possible -- at some additional engineering and hardware cost -- to design acyclic circuits of lenient components that are themselves lenient combinational devices.

In general, non-lenience of combinational circuits stems from the dependence of an output on any of several input-output propagation paths. In the case of our 3-input NAND-NAND circuit, observe that for $C=1$ the path holding the output $Y=1$ involves circuit nodes $C=1 \rightarrow R=0 \rightarrow Y=1$; when $C=0$, the alternative path $C=0 \rightarrow P=1 \rightarrow Q=0 \rightarrow Y=1$ forces the output to be 1. Following the $1 \rightarrow 0$ transition of C , the former path is disabled (no longer forcing $Y=1$) and the latter path is enabled (forcing $Y=1$), but the timing of these two events is not strictly controlled. Hence there is a period during which neither path may be active, and the value of Y is hence undetermined.

Using the terminology of Section 7.6, our circuit computes the value $Y=\bar{C}A+CB$ via two implicants, CA and $\bar{C}B$, computed by separate paths and combined in the final logical OR (using a Demorgitized NAND). The C transistion switches responsibility for holding $Y=1$ from one path to the other, effectivly selecting whether the A input or the B input should be taken as the output value Y . In our test case, $A=B=1$ rendering the output value independent of which path is selected; however, there is no single path in our circuit which forces $Y=1$ when $A=B=1$.

We can address this problem by adding a path representing a third implicant, AB , to the circuit. The circuit shown to the right uses three NAND gates to compute inverted versions of the three implicants $A\bar{C}$, AB , and CB , and the final 3-input NAND computes their sum. The resulting value is $Y=A\bar{C}+AB+CB$, which is logically equivalent to our simplified $Y=A\bar{C}+CB$; however, now there is a path dedicated to ensuring $Y=1$ for each of three (rather than two) implicants.

Lenient NAND-NAND Circuit

This approach can be generalized to the implementation of lenient combinational circuits to compute arbitrary Boolean expressions. In the general case, a sum-of-products implementation using lenient components will be lenient if it includes logic representing every prime implicant of the function.

Fan-In and Delay

7.7.2. Fan-In and Delay

Although the CMOS gate examples have focused on 2-input NAND and NOR gates, we can of course implement many functions with more inputs -- higher fan-in -- as single CMOS gates. For example, the diagram on the right shows the straightforward extension of our CMOS NAND gate to 4 inputs, by including additional PFETs and NFETs gated by the new inputs to the parallel and series chains of the pullup and pulldown circuits. While this approach can in theory be extended to single CMOS gates of arbitrarily many inputs, the performance of CMOS gates degrades with increasing fan-in due to electrical loading. In practice, it is typically preferable to implement functions of more than about four inputs as circuits of smaller gates.

4-input NAND gate

Certain common logical functions like AND, OR, and XOR offer a variety of alternative approaches to large-fan-in circuit design. Each of these functions F have the property that $F(a_1,\dots,a_k)=F(a_1,F(a_2,\dots,a_k))$, suggesting a strategy for repeatedly using 2-input gates to reduce the fan-in requirement by one. But we may also observe that for each such function F , $F(a_1,\dots,a_k)=F(F(a_1,\dots,a_j),F(a_{j+1},\dots,a_k))$: we may partition the k arguments into multiple sub-operations, and employ a divide-and-conquer approach to their combination.

2-input XOR

Suppose, for example, we have a 2-input XOR gate shown on the left. Note from the truth table that the XOR function has the value 1 if and only if an odd number of its inputs are 1, a function that extends in a straightforward way to arbitrarily many inputs and that has the properties described for F above. We desire to implement $XOR(a_1,\dots,a_N)$

for arbitrary fan-in, a function that returns 1 if the number of 1s among its N inputs is odd.

This function, sometimes called odd parity, is useful for error detection techniques such as those discussed in Section 3.4.

The observation that $XOR(a_1, \dots, a_k) = XOR(a_1, XOR(a_2, \dots, a_k))$ suggests that we can compute $XOR(a_1, \dots, a_N)$ by the circuit shown to the right, in which each of the two-input XOR gates combines a single input with the cumulative XORs from the gates to its left. While this circuit works correctly, we note that the longest input-output path goes through all $N-1$ of the component 2-input gates, yielding a propagation delay that grows in proportion to the number of inputs.

While this may be an acceptable approach in some circumstances, we often prefer an approach that minimizes the worst-case propagation delay of the resulting circuit. For associative operations like XOR , we observe that $XOR(a_1, \dots, a_k) = XOR(XOR(a_1, \dots, a_j), XOR(a_{j+1}, \dots, a_k))$ suggesting that we can reduce our N -input XOR to the XOR of the result of two $N/2$ -input XORs -- a step that halves the fan-in at the cost of a single gate delay.

Repeated application of this divide-and-conquer step results in a binary tree structure as shown to the left. Note that a tree having k levels of 2-input components has a propagation delay proportional to k , but provides a fan-in of 2^k inputs. Thus an N -input XOR, for large fan-in N can be computed by the tree in time proportional to $\log_2(N)$, an improvement over the linear delay of the preceding approach.

Similar tree structures can be used to compute many other functions whose extension to N arguments is straightforward, including AND, OR, and addition. Certain functions (like NAND and NOR) admit tree implementations, but require some care due to their special properties (e.g., inversion). Other functions don't decompose so nicely, and don't lend themselves to tree implementations; an example is the N -input majority function, which returns 1 if a majority of its inputs are 1.

Systematic Implementation Techniques

7.8. Systematic Implementation Techniques

Typically the smallest, fastest, and most cost-effective implementation of a complex Boolean function is a custom circuit designed and optimized to compute that function. While there are many sophisticated techniques and tools for designing such circuits, the customization process is inherently expensive -- both in engineering and manufacturing costs. For this reason, several synthesis approaches have been developed that forego function-specific optimizations in order to allow the straightforward implementation of an arbitrary k -input Boolean function by parameterizing a fixed, function-generic hardware substrate. In this section we sample several such approaches.

Multiplexors

7.8.1. Multiplexors

The three-input Boolean function used as a running example in this Chapter is a a commonly-used component in our combinational repertoire: a 2-way multiplexor (or Mux in the vernacular of logic design). The logic symbol and truth table for a 2-way mux is shown to the right. A useful model for its behavior is a 2-way logic-controlled switch: it takes 2 independent data inputs D_0 and D_1 , as well as a select input S ; its output (after a propagation delay) is D_0 if $S=0$ or D_1 if $S=1$. Muxes are useful for routing the flow

of data through our combinational circuits on the basis of control signals applied to their select inputs; we will see many examples of such uses in subsequent chapters. As discussed in previous sections of this Chapter, the implementation of a multiplexor may or may not be lenient; the example of section Section 7.7.1 is, in fact, a lenient 2-way multiplexor. This device plays a key role in the implementation of memory devices in the next Chapter.

The multiplexor function can be generalized to 2^k inputs and k select lines; the diagram to the left shows a 4-way multiplexor, whose output is one of the four data inputs as selected by the four possible combinations of inputs on the two select lines. One way to implement a 2^k -way multiplexor is to use a

4-way Multiplexor

2-way Multiplexor

2-way mux selecting among the output of two 2^{k-1} -way muxes, as illustrated in the 4-way mux implementation shown to the right.
Lookup Table Synthesis

7.8.2. Lookup Table Synthesis

The ability of a multiplexor to select one of 2^k input values suggests its use as a mechanism to perform table lookup functions: based on its select input, its output is one of the connected data values. We might implement an arbitrary 2-input Boolean function, for example, by use of a 4-way mux as shown to the left. The Boolean inputs drive the select lines, and hence select one of the four data inputs. If we connect the data inputs to constant 0 and 1 values corresponding to the output column of the Boolean function's truth table, the supplied select inputs will effectively look up the correct output from the supplied values. The function "programmed" into this implementation is NAND, but by simply changing the constant values -- 0s and 1s implemented via V_{DD} and ground connections -- we can make this circuit compute any of the 16 2-input Boolean functions.

Using larger multiplexors, we can implement functions with more inputs. The diagram to the right shows the implementation of a 3-input function -- the carry output for a 1-bit full adder device -- using an 8-way multiplexor. Again, the mux simply chooses among the constant values supplied (in accordance with a truth table) as its data inputs.

Table lookup approaches are attractive in applications where it is desirable to isolate the hardware structure from the choice of the function to be implemented, for example when that function might be frequently changed or is specified as a customization parameter of an otherwise commodity hardware device.

Read-only Memories

7.8.3. Read-only Memories

When the number of inputs becomes large, the use of a simple mux-based approach to table lookup is tedious: a 10-input function, for example, requires a 1024-way multiplexor. For this reason, table lookup approaches have evolved to space-efficient structures that more effectively exploit the 2-dimensional layout constraints of the surface of a silicon chip. Such approaches are often called read-only memories (ROMs), and share with the mux-based lookup table the fact that the output values are programmed into the structure as data values in locations corresponding to combinations of inputs. The various input combinations may be viewed as addresses of memory locations, and the data values as the memory contents.

This section provides a glimpse of a simple read-only memory technology for the systematic implementation of arbitrary Boolean functions. Although the technology is compatible with CMOS logic, it is noteworthy that it differs from CMOS gate implementation in a number of important ways -- among them are lenience and power consumption characteristics.

Decoders

7.8.3.1. Decoders

In order to expand our table lookup implementation to two dimensions, we introduce another building block: a 2^k -way decoder. The circuit symbol for a decoder is shown to the right. It takes as input k select lines, whose values select one of 2^k output lines to be driven to a logical 1; all unselected lines remain at a logical 0. Thus for any combination of valid inputs, one and only one of the output lines will be 1: the output lines may be viewed a decoded (or "one-hot") representation of the binary value supplied on the k select inputs.

ROM-based Full Adder

7.8.3.2. ROM-based Full Adder

Consider the simple 3-input, 2-output combinational device whose symbol and truth table are shown to the left. The device is in fact a 1-bit full adder, a building block used in the construction of binary adders in subsequent chapters; for now, we use it simply as an example of combinational functions we might wish to implement.

Full Adder

to the left. The device is in fact a 1-bit full adder, a building block used in the construction of binary adders in subsequent chapters; for now, we use it simply as an example of combinational functions we might wish to implement.

The device takes three binary inputs A , B , and C_i and produces the two outputs S and C_o . Note that we could have split this into two separate 3-input devices for computing S and C_o respectively; however, their combination as a single device allows us to exploit some shared logic for implementation efficiency. In particular, we will use a single decoder to convert the three input bits to a more convenient form: a logic 1 on exactly one of its eight outputs.

One approach to the ROM-based full adder is shown to the right. The three inputs are decoded to drive exactly one of the eight horizontal lines to 1; the selected line corresponds to the row of the truth table selected by the current inputs, or one possible product term in a sum-of-products expression.

ROM-based Full Adder

Superimposed over the horizontal decoder outputs are two vertical lines, each connected to V_{DD} through a resistive pullup. This passive pullup serves to cause the voltage on each vertical line to have a default value of 1. At each of the intersections between a horizontal decoder output and a vertical line, there is either no connection or an NFET pulldown. The presence of a pulldown at an intersection causes a 1 on the horizontal line to force the vertical line to 0, by overpowering the passive resistive pullup. Thus each of the vertical lines performs a logical NOR function of the horizontal lines for which there are pulldowns; an inverter between each vertical line and the S and C_o outputs causes each output to be a logical OR of the decoder outputs selected by the presence of pulldowns.

Comparing the diagram with the truth table for the full adder, note that there are pulldowns at precisely those intersections corresponding to 1s in the truth table output columns. Thus the logical function that appears at each output corresponds to the sum-of-products computation of the functions represented in the table.

ROM Folding

7.8.3.3. ROM Folding

Although this ROM implementation exploits a second dimension to compute multiple outputs from a single decoded set of inputs, its vertical dimension still grows in proportion to 2^k where k is the number of inputs. In general, long lines cause delays; we would prefer to make the ROM array more nearly square by using a greater number of shorter vertical lines.

Folded ROM

Fortunately, this is a fairly straightforward engineering tradeoff. On the left is shown a revised ROM layout, in which two of the inputs (A and B) are decoded onto four horizontal lines rather than the previous eight. Each output is then computed by two vertical lines, one giving the value for $C_i = 0$ and another for $C_i = 1$; the C_i input is used in an output multiplexor for each output to select the proper value.

Commodity/Customization Tradeoffs

7.9. Commodity/Customization Tradeoffs

The design approach taken by ROMs, involving a regular array of cells sharing the same internal structure, is used in other systematic logic technologies as well. These include read-write memories (RAMs) as well as several varieties of logic arrays.

These approaches exemplify a recurring technology tradeoff between the advantages of an application-generic, mass-produced commodity implementation and an application-specific implementation highly optimized around the structure of a particular solution. An attractive feature of ROMs and other commoditized approaches is that nearly all of the structure of a given-size ROM is independent of the actual function it performs: the customization may be left to a final parameterization step (adding a metal layer to a chip, or even field programming in the case of flash memories or FPGA logic arrays).

7.10. Chapter Summary

Having moved from transistors into the logic domain, we face the task of combining logic elements to synthesize higher-level logical functions. Techniques explored include:

- While truth tables are an effective tool for the specification of logic functions, Boolean expressions are often a convenient alternative. Familiarity with identities and Boolean manipulation techniques is an important skill for the logic designer.
- Any truth table can easily be translated to a sum-of-products (SOP) Boolean expression, which can be mechanically converted to a 3-level gate implementation of the function. Useful SOP techniques include NAND-NAND and NOR-NOR implementations.
- SOP implementations are not necessarily minimal; function-specific structure may often be exploited to implement a given function using fewer CMOS devices.
- Multiplexors are useful building blocks, and can be used to systematically implement arbitrary functions as lookup tables.
- ROMs generalize this table-lookup implementation approach, offering a fixed hardware complexity independently of the truth table details.